

# Getting Started with PostgreSQL

M. Edward (Ed) Borasky

2018-03-02

# Introduction: PostgreSQL in Context

## Industry trends: NoSQL databases (Redmond and Wilson 2012)

- ▶ Graph databases
- ▶ Key-value stores
- ▶ JSON document stores
- ▶ In-memory databases

## Industry trends: Backend frameworks (Apache/PHP, Ruby on Rails, Django)

- ▶ Work with *any* database (MySQL/MariaDB, SQLite, PostgreSQL)
- ▶ Just use the database for CRUD, application logic is all in PHP / Ruby / Python / JavaScript code!
- ▶ Wait - CRUD?
  - ▶ Create
  - ▶ Read (aka SELECT)
  - ▶ Update
  - ▶ Delete

## If all you want . . .

- ▶ If all you want is an industrial-strength open-source permissive-licensed CRUD engine that
  - ▶ Is fully ACID compliant
  - ▶ Scales to huge installations
  - ▶ Has replication / failover / high availability as standard equipment
- ▶ Yeah, PostgreSQL's got that.

## But if you also want . . .

- ▶ Full-text search
- ▶ Stored procedures in Python, Perl, Ruby, R, Tcl and Lua
- ▶ Foreign data (Text files, GIS data, MySQL/MariaDB, Redis) mapped into your database (foreign data wrappers)
- ▶ Key-value stores (hstore)
- ▶ JSON document stores (jsonb)
- ▶ Yeah, PostgreSQL's got that too!

## Speaking of industrial strength ...

- ▶ Geographic Information Systems (GIS)
- ▶ PostGIS
  - ▶ Read and write GIS data files
  - ▶ Process geometric, geographic and topology GIS data types
  - ▶ Both vector and raster data
  - ▶ Geocoding, reverse geocoding, address standardization
- ▶ pgRouting
  - ▶ Shortest / fastest / lowest cost routes from point A to point B
  - ▶ Traveling salesperson problem
  - ▶ Turn-by-turn directions for cars, bikes and pedestrians!
- ▶ Yeah, I want that CRUD too!

## PostgreSQL on the Desktop - single user



## Windows or Mac

- ▶ Go to EnterpriseDB download site  
<https://www.enterprisedb.com/downloads/postgres-postgresql-downloads>.
- ▶ Select the latest version (10.3)
- ▶ Select your operating system (Windows or Mac)
- ▶ Installation: install everything but don't run StackBuilder yet

## Linux: use the PGDG repositories

- ▶ RHEL / CentOS / Fedora:  
<https://www.postgresql.org/download/linux/redhat/>
- ▶ Ubuntu: <https://www.postgresql.org/download/linux/ubuntu/>  
(probably works for Linux Mint)
- ▶ Debian: <https://www.postgresql.org/download/linux/debian/>
- ▶ Installation: install PostgreSQL 10.3 and pgAdmin 4 for desktop

## Connecting with pgAdmin

- ▶ Reference: R. Obe and Hsu (2017b), chapter 4

## Exploring the tree

## Creating a database

## Creating tables - Data Definition Language (DDL)

To create a table, you need to do two things:

1. Define the names and data types of every column in the table.
2. Load the data into the table.

We'll be working with a sample file of Oregon highway mileposts.

You'll find it in

[https://github.com/hackoregon/data-science-pet-containers/blob/master/examples/mileposts/Mileposts\\_2014/Mileposts\\_2014.csv](https://github.com/hackoregon/data-science-pet-containers/blob/master/examples/mileposts/Mileposts_2014/Mileposts_2014.csv).

## A cheat code if you're in a hurry

If you don't know the data types, you can always just set them all to text and re-cast them to the correct type later! But in this case it's mostly obvious which columns are numeric or timestamps, and we can use text for the rest.

For the mileposts, we'll use `double precision` for the latitude and longitude and `date` for `GIS_PRC_DATE` and `text` for the others.

## The finished DDL

```
CREATE TABLE mileposts_2014 (  
    hwyname text, hwynumb text, st_hwy_sfx text,  
    rdwy_id text, mlge_typ text, ovlp_cd text,  
    mp text, mp_desc text, mp_disp text,  
    lrs_key text, lrm_key text,  
    lat double precision, longtd double precision,  
    hrz_col_m text, crd_rf_dtm text, effectv_dt text,  
    gis_prc_dt date);
```



## Hey, Ed, what's with all the lower case names?

PostgreSQL requires special care in coding SQL queries when column names have anything besides lower-case letters, numbers or underscores. You have to enclose them in double-quotes. It's a real hassle, so “snake\_case” rules!

## Reading a CSV file into a table - COPY

```
COPY mileposts_2014 FROM  
'/d/Projects/data-science-pet-containers/examples/mileposts  
WITH CSV HEADER;
```

## Backing up a database

## Restoring a database

## Creating a query

## PostgreSQL on a (Linux) Server

## Two dialects of Linux

- ▶ Ubuntu 16.04.x LTS (most popular)
- ▶ RHEL 7 / CentOS 7

## Docker container

- ▶ Debian stable - similar to Ubuntu
- ▶ Alpine - avoid this unless you want to do a lot of research
- ▶ I have a full PostgreSQL / PostGIS / pgRouting stack in a container at

<https://github.com/hackoregon/data-science-pet-containers>.

- ▶ See <https://github.com/hackoregon/data-science-pet-containers/blob/master/containers/small.yml> for the Docker compose file
- ▶ The documentation's a bit sparse still
- ▶ Native executables will perform better
- ▶ Docker hosting on a desktop / laptop isn't exactly end-user friendly yet



## The PGDG repositories

- ▶ Linux distributions ship with the version of PostgreSQL that was stable when their release process froze features.
- ▶ You'll still get security updates and bug fixes, but no new features.
- ▶ So the PostgreSQL Global Development Group (PGDG) maintains up-to-date binary repositories for all the major Linux distros.
- ▶ These binaries are as well tested and supported as those that ship with the Linux distros.
- ▶ Download page: <https://www.postgresql.org/download/>

## Configuring

- ▶ Initializing the cluster
- ▶ Enabling and starting the service

## Listening address and port

# Authentication

## Fine-grained permissions and roles

# Tablespaces

# PostGIS

## Reading in GIS data (Obe and Hsu 2015, chap. 4)

- ▶ Shapefiles
- ▶ “GDB” databases
- ▶ OpenStreetMap data



## Tagging points with a geometry column

## Geocoding (Obe and Hsu 2015, chap. 8)

- ▶ There's a demo using a Docker container at <https://github.com/hackoregon/data-science-pet-containers/tree/master/examples/geocoding>.
- ▶ This uses a container running Debian Linux with a full PostgreSQL / PostGIS / pgRouting stack.
- ▶ ***I have not tested it on a Windows Docker host, just Linux. It undoubtedly needs some adjustments for the volumes mounted on host filesystems.***

# Spatial joins

## pgRouting - what it can do

## Turn-by-turn directions

- ▶ Reference: (R. Obe and Hsu 2017a)

# References

## References

Obe, Regina O., and Leo S. Hsu. 2015. *PostGIS in Action, Second Edition*. Manning Publications Co. <https://www.manning.com/books/postgis-in-action-second-edition>.

Obe, Regina, and Leo Hsu. 2017a. *PgRouting: A Practical Guide*. Locate Press LLC. <https://locatepress.com/pgrouting>.

———. 2017b. *PostgreSQL: Up and Running, 3rd Edition*. O'Reilly Media, Inc. <http://shop.oreilly.com/product/0636920052715.do>.

Redmond, Eric, and Jim R. Wilson. 2012. *Seven Databases in Seven Weeks*. The Pragmatic Programmers. <https://pragprog.com/book/rwdata/seven-databases-in-seven-weeks>.