

Getting Started with PostgreSQL

M. Edward (Ed) Borasky

2018-03-02

Introduction: PostgreSQL in Context

Industry trends: NoSQL databases (Redmond and Wilson 2012)

- ▶ Graph databases
- ▶ Key-value stores
- ▶ JSON document stores
- ▶ In-memory databases

Industry trends: Backend frameworks (Apache/PHP, Ruby on Rails, Django)

- ▶ Work with *any* database (MySQL/MariaDB, SQLite, PostgreSQL)
- ▶ Just use the database for CRUD, application logic is all in PHP / Ruby / Python / JavaScript code!
- ▶ Wait - CRUD?
 - ▶ Create
 - ▶ Read (aka SELECT)
 - ▶ Update
 - ▶ Delete

If all you want . . .

- ▶ If all you want is an industrial-strength open-source permissive-licensed CRUD engine that
 - ▶ Is fully ACID compliant
 - ▶ Scales to huge installations
 - ▶ Has replication / failover / high availability as standard equipment
- ▶ Yeah, PostgreSQL's got that.

But if you also want . . .

- ▶ Full-text search
- ▶ Stored procedures in Python, Perl, Ruby, R, Tcl and Lua
- ▶ Foreign data (Text files, GIS data, MySQL/MariaDB, Redis) mapped into your database (foreign data wrappers)
- ▶ Key-value stores (hstore)
- ▶ JSON document stores (jsonb)
- ▶ Yeah, PostgreSQL's got that too!

Speaking of industrial strength ...

- ▶ Geographic Information Systems (GIS)
- ▶ PostGIS
 - ▶ Read and write GIS data files
 - ▶ Process geometric, geographic and topology GIS data types
 - ▶ Both vector and raster data
 - ▶ Geocoding, reverse geocoding, address standardization
- ▶ pgRouting
 - ▶ Shortest / fastest / lowest cost routes from point A to point B
 - ▶ Traveling salesperson problem
 - ▶ Turn-by-turn directions for cars, bikes and pedestrians!
- ▶ Yeah, I want that CRUD too!

PostgreSQL on the Desktop - single user

Windows or Mac

- ▶ Go to EnterpriseDB download site
<https://www.enterprisedb.com/downloads/postgres-postgresql-downloads>.
- ▶ Select the latest version (10.3)
- ▶ Select your operating system (Windows or Mac)
- ▶ Installation: install everything but don't run StackBuilder yet

Linux: use the PGDG repositories

- ▶ RHEL / CentOS / Fedora:
<https://www.postgresql.org/download/linux/redhat/>
- ▶ Linux Mint 18 / Ubuntu 16.06 LTS ("Xenial"):
<https://github.com/hackoregon/data-science-pet-containers/blob/master/docs/Mint18-Xenial/README.md>
- ▶ Debian: <https://www.postgresql.org/download/linux/debian/>
- ▶ Installation: install PostgreSQL 10.3 and pgAdmin 4 for desktop

Connecting with pgAdmin - (R. Obe and Hsu (2017b), chapter 4)

1. Right-click on “Servers”
2. Create a server
3. Give it any name you want
4. Fill in the connection tab

EnterpriseDB connection

- ▶ Host: localhost
- ▶ Port: 5432
- ▶ Maintenance database and user are postgres
- ▶ Password is what you set when you installed

Linux connection (at least on Mint / Ubuntu)

- ▶ Host: /var/run/postgresql
- ▶ Port: 5432
- ▶ Maintenance database and user are the same as your Linux ID
- ▶ Password is empty

Exploring the tree (pgAdmin live)

Creating a database (pgAdmin live)

Creating tables - Data Definition Language (DDL)

To create a table, you need to do two things:

1. Define the names and data types of every column in the table.
2. Load the data into the table.

We'll be working with a sample file of Oregon highway mileposts.
You'll find it in

https://github.com/hackoregon/data-science-pet-containers/blob/master/examples/mileposts/Mileposts_2014/Mileposts_2014.csv.

A cheat code if you're in a hurry

If you don't know the data types, you can always just set them all to `text` and re-cast them to the correct type with SQL later! But in this case it's mostly obvious which columns are numeric or timestamps, and we can use `text` for the rest.

For the mileposts, we'll use `double precision` for the latitude and longitude and `text` for the others.

The DDL

```
CREATE TABLE mileposts_2014 (  
    hwyname text, hwynumb text, st_hwy_sfx text,  
    rdwy_id text, mlge_typ text, ovlp_cd text,  
    mp text, mp_desc text, mp_disp text,  
    lrs_key text, lrm_key text,  
    lat double precision, longtd double precision,  
    hrz_col_m text, crd_rf_dtm text, effectv_dt text,  
    gis_prc_dt text);
```

Hey, Ed, what's with all the lower case names?

PostgreSQL requires special care in coding SQL queries when column names have anything besides lower-case letters, numbers or underscores. You have to enclose them in double-quotes. It's a real hassle, so “snake_case” rules!

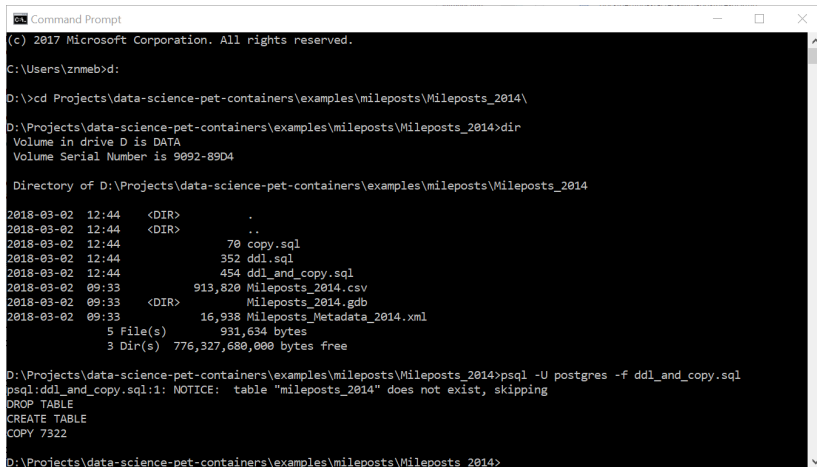
Loading the data into the table: `psql \copy!`

- ▶ `psql` is the command-line PostgreSQL client.
- ▶ It's fully scriptable - you can write programs in it.
- ▶ It has string substitution.
- ▶ You can mix SQL statements and `psql` commands
- ▶ It's great for reproducibility
- ▶ But it's another language to learn
 - ▶ most folks will script in a language they know, like Python, which has its own PostgreSQL client libraries.

Mixing SQL and psql

- ▶ Script: https://github.com/hackoregon/data-science-pet-containers/blob/master/examples/mileposts/Mileposts_2014/ddl_and_copy.sql
- ▶ To run it, we type `psql -f ddl_and_copy.sql`.

This is what it looks like in a Windows cmd window



```
Command Prompt
(c) 2017 Microsoft Corporation. All rights reserved.

C:\Users\znmeb>d:

D:\>cd Projects\data-science-pet-containers\examples\mileposts\Mileposts_2014\

D:\Projects\data-science-pet-containers\examples\mileposts\Mileposts_2014>dir
Volume in drive D is DATA
Volume Serial Number is 9092-89D4

Directory of D:\Projects\data-science-pet-containers\examples\mileposts\Mileposts_2014

2018-03-02  12:44    <DIR>          .
2018-03-02  12:44    <DIR>          ..
2018-03-02  12:44             70 copy.sql
2018-03-02  12:44            352 ddl.sql
2018-03-02  12:44           454 ddl_and_copy.sql
2018-03-02  09:33        913,820 Mileposts_2014.csv
2018-03-02  09:33    <DIR>        Mileposts_2014.gdb
2018-03-02  09:33       16,938 Mileposts_Metadata_2014.xml
                5 File(s)          931,634 bytes
                3 Dir(s)  776,327,680,000 bytes free

D:\Projects\data-science-pet-containers\examples\mileposts\Mileposts_2014>psql -U postgres -f ddl_and_copy.sql
psql:ddl_and_copy.sql:1: NOTICE:  table "mileposts_2014" does not exist, skipping
DROP TABLE
CREATE TABLE
COPY 7322

D:\Projects\data-science-pet-containers\examples\mileposts\Mileposts_2014>
```

```

DROP TABLE IF EXISTS mileposts_2014;
CREATE TABLE mileposts_2014 (
    hwyname text,
    hwynumb text,
    st_hwy_sfx text,
    rdwy_id text,
    mlge_typ text,
    ovlp_cd text,
    mp text,
    mp_desc text,
    mp_disp text,
    lrs_key text,
    lrm_key text,
    lat double precision,
    longtd double precision,
    hrz_col_m text,
    crd_rf_dtm text,
    effectv_dt text,
    gis_prc_dt text
);
\copy mileposts_2014 from 'Mileposts_2014.csv' with csv header

```

Examining the table (pgAdmin live)

Backing up a database: `pg_dump` (pgAdmin live)

Restoring a database: `pg_restore` (pgAdmin live)

PostgreSQL on a (Linux) Server

And now for something completely different

- ▶ On a desktop, you're usually the database superuser
 - ▶ Everything works, and your machine's behind a firewall
- ▶ On a server, there's a DBA and a SysAdmin that hands out privileges only with justification
 - ▶ You don't get it unless you need it, and everything you do is logged and monitored
- ▶ Bottom line: don't be surprised if you get told, "No! Find another way."

Two most common server dialects of Linux

- ▶ Ubuntu 16.04.x LTS (most popular)
- ▶ RHEL 7 / CentOS 7

An aside - PostgreSQL in a Docker container

- ▶ Usually Debian stable - similar to Ubuntu
- ▶ Sometimes Alpine - avoid this unless you want to do a lot of research
- ▶ I have a full PostgreSQL / PostGIS / pgRouting stack in a container at

<https://github.com/hackoregon/data-science-pet-containers>.

- ▶ See <https://github.com/hackoregon/data-science-pet-containers/blob/master/containers/small.yml> for the Docker compose file
- ▶ The documentation's a bit sparse still
- ▶ Native executables will perform better
- ▶ Docker hosting on a desktop / laptop isn't exactly end-user friendly yet

The PGDG repositories

- ▶ Linux distributions ship with the version of PostgreSQL that was stable when their release process froze features.
- ▶ You'll still get security updates and bug fixes, but no new features.
- ▶ So the PostgreSQL Global Development Group (PGDG) maintains up-to-date binary repositories for all the major Linux distros.
- ▶ These binaries are as well tested and supported as those that ship with the Linux distros.
- ▶ Download page: <https://www.postgresql.org/download/>

Configuring

- ▶ Initializing the cluster: `initdb`
- ▶ Enabling and starting the service: `systemctl`

Listening address and port

- ▶ Edit `postgresql.conf`
- ▶ Restart the service

Authentication

- ▶ Edit `pg_hba.conf`
- ▶ Restart the service

Fine-grained permissions and roles

- ▶ Adding a user (pgAdmin live)
- ▶ SQL GRANT and REVOKE

PostGIS

Installing PostGIS

- ▶ Windows and Mac - use the Stack Builder
 - ▶ You'll get `pgRouting` too
 - ▶ You'll get the command-line GIS utility `gdal`
- ▶ Linux: use the PGDG repositories
 - ▶ You'll need to install `pgRouting` and `gdal` explicitly

Reading in GIS data (Obe and Hsu 2015, chap. 4)

- ▶ Shapefiles (live command line demo)
- ▶ “GDB” databases (live command line demo)
- ▶ OpenStreetMap data (live command line demo)

Tagging points with a geometry column (pgAdmin live)

Geocoding (Obe and Hsu 2015, chap. 8)

- ▶ There's a demo using a Docker container at <https://github.com/hackoregon/data-science-pet-containers/tree/master/examples/geocoding>.
- ▶ This uses a container running Debian Linux with a full PostgreSQL / PostGIS / pgRouting stack.
- ▶ ***I have not tested it on a Windows Docker host, just Linux. It undoubtedly needs some adjustments for the volumes mounted on host filesystems.***

pgRouting - what it can do

Turn-by-turn directions

- ▶ Reference: (R. Obe and Hsu 2017a)

References

References

Obe, Regina O., and Leo S. Hsu. 2015. *PostGIS in Action, Second Edition*. Manning Publications Co. <https://www.manning.com/books/postgis-in-action-second-edition>.

Obe, Regina, and Leo Hsu. 2017a. *PgRouting: A Practical Guide*. Locate Press LLC. <https://locatepress.com/pgrouting>.

———. 2017b. *PostgreSQL: Up and Running, 3rd Edition*. O'Reilly Media, Inc. <http://shop.oreilly.com/product/0636920052715.do>.

Redmond, Eric, and Jim R. Wilson. 2012. *Seven Databases in Seven Weeks*. The Pragmatic Programmers. <https://pragprog.com/book/rwdata/seven-databases-in-seven-weeks>.